

Hour 3 – Detailed Component Design & Technology Choices

Goal: Provide concrete technology selections, detailed component designs, example schemas, API contracts, sequence diagrams (text), deployment patterns, and differential privacy mechanics to convert the Hour-2 architecture into an implementable blueprint.

0) Design Principles (recap)

- Keep the system **modular**: clear separation of ingress, processing, policy, storage, and governance.
 - Prefer **battle-tested OSS** with cloud-native managed options for operational simplicity.
 - Optimize for **observability, security, and deterministic latency**.
-

1) Technology Stack (recommended)

- **Ingress & Messaging**
 - API Gateway: **Kong / Envoy / AWS API Gateway (regional)**
 - Messaging backbone: **Apache Kafka (Confluent)** or **Redpanda** (single-binary, lower ops)
 - Lightweight pub/sub for control plane: **NATS JetStream**
- **Compute & Orchestration**
 - Kubernetes (EKS/GKE/AKS or on-prem k8s)
 - Container runtime: containerd
 - Service mesh: **Istio** or **Linkerd** for mTLS & telemetry
- **Storage**
 - Object store: **S3 / MinIO** (for on-prem)
 - Columnar lake: **Delta Lake / Apache Iceberg** on top of S3
 - Metadata store: **Postgres** (timescaled for audit metadata if needed)
 - Ledger: **Trillian / immudb** for append-only proofs
 - Search: **OpenSearch** (for PHI-free indices)
- **Inference & ML**
 - Model serving: **NVIDIA Triton, TorchServe** (fallback)
 - Model training infra: **KubeFlow, MLflow** for registry
 - FL coordinator: **Flower** / custom coordinator with gRPC

- DP libs: **TensorFlow Privacy**, **Opacus** (PyTorch), custom moments accountant
- **OCR/ASR/Imaging**
 - OCR: **Tesseract** (baseline) + commercial options (Google/Vision or AWS Textract) or custom Vision-LM models
 - ASR: **WhisperX**, **Kaldi** pipelines, or Whisper running on GPU for accuracy + diarization (pyannote, ECAPA)
 - DICOM: **pydicom**, **Orthanc** for PACS integration
- **Policy & Authorization**
 - Policy engine: **Open Policy Agent (OPA)** or **Cedar**
 - AuthN/AuthZ: **Keycloak** / **Auth0** or cloud IAM
- **KMS/HSM & Crypto**
 - Cloud KMS (AWS KMS, Google KMS) + HSM (Cloud HSM or Thales) for key escrow
 - Envelope encryption libraries: **libsodium**, **Google Tink**
- **Observability**
 - Metrics: **Prometheus**
 - Tracing: **Jaeger** or **OpenTelemetry**
 - Logs: **Loki** + Grafana
 - Dashboarding: **Grafana** (SLO dashboards)
- **CI/CD & Security**
 - GitHub Actions / GitLab CI
 - Image scanning: **Clair**, **Trivy**
 - SAST/DAST: **SonarQube**, **OWASP ZAP**

2) Component Designs & Interfaces

2.1 Ingestion API (Concrete Contract)

Endpoint: POST /v1/documents **Headers:** Authorization: Bearer <JWT>, X-Tenant-Id, Content-Type: multipart/form-data **Query params:** use_case=[billing|research|clinical|analytics]

Body (multipart) - `metadata` (JSON): { "document_id": "uuid", "source_system": "epic", "received_at": "2025-08-28T10:12:00Z", "format": "pdf" } - `file` (binary)

Response (sync small docs) - 200 OK
{ "document_id":"uuid","status":"accepted","ingest_token":"abc123" }

Semantic guarantees: idempotent ingest via `document_id` + `X-Tenant-Id`.

2.2 Redaction API (Sync with SLA)

Endpoint: `POST /v1/redact` **Body:**

```
{
  "document_id":"uuid",
  "tenant":"acme_hosp",
  "use_case":"research",
  "content":"base64...",
  "content_type":"text/plain"
}
```

Response: { "redacted_content":"base64...", "pseudo_ids":{"..."}, "audit_hash":"hex" }

SLA: Small docs (<32KB) processed sync with p95<100ms.

2.3 Event Schema (Canonical Avro Example)

```
{
  "type": "record",
  "name": "ClinicalDocument",
  "namespace": "com.company.health",
  "fields": [
    { "name": "document_id", "type": "string" },
    { "name": "tenant_id", "type": "string" },
    { "name": "ingest_ts", "type": { "type": "long", "logicalType": "timestamp-millis" } },
    { "name": "modality", "type": "string" },
    { "name": "payload_ref", "type": "string" },
    { "name": "pii_spans", "type": [ "null", { "type": "array", "items": {
      "type": "record", "name": "PiiSpan", "fields": [
        { "name": "type", "type": "string" },
        { "name": "start", "type": "int" },
        { "name": "end", "type": "int" },
        { "name": "confidence", "type": "float" }
      ]
    } ] } ]
  ]
}
```

```

    ]
  }}, "default": null}
]
}

```

3) Sequence Diagrams (Textual)

3.1 Ingest → OCR → NER → Redact (sync)

1. Client -> API Gateway: POST /v1/documents
2. API Gateway -> Ingest Service: validate, stamp tenant, push to Kafka raw.documents
3. OCR Worker consumes raw.documents, emits normalized.documents
4. NER Worker consumes normalized.documents, emits pii_found
5. Policy Engine (sync call) returns redaction recipe
6. Redactor applies FPE / pseudonymization → writes redacted.documents
7. Audit Service hashes input+policy+output → appends to ledger
8. Ingest Service returns sync 200 with audit_hash

3.2 Audio Streaming (Realtime)

1. Client -> WebSocket/gRPC Stream opens
2. Chunk producer -> Stream Ingest service (Kafka streaming or NATS)
3. ASR DAG processes frames, emits partial transcripts
4. Diarization attaches speaker IDs; NER marks PHI spans with timestamps
5. Redactor masks live transcript and emits sanitized stream to subscriber
6. Audit events buffered & anchored periodically

4) Pseudonymization Design

- **Stable Pseudo-IDs:** pseudo_id = HMAC_SHA256(salt || canonical_entity_key) where salt = per-tenant secret stored in KMS and canonical_entity_key = normalized name+DOB+hash(MPI_hint).
- **Entity Graph:** Graph DB (Neo4j / DGraph) to maintain relationships (family links) and cross-reference identities; mapping stored as entity nodes with deterministic pseudo_id.
- **Reversal policy:** De-pseudonymization requires multi-party approval; keys accessed only via HSM with audit gating.

5) Differential Privacy: Operational Mechanics

- **Per-request DP Mode:** API caller may request DP sanitation with dp_profile header. The policy engine reserves ϵ from tenant budget.

- **Mechanisms:** DP-SGD for model releases; Laplace/Gaussian mechanisms for numeric aggregates; randomized response for categorical where applicable.
 - **Accounting:** Moments accountant implementation (reference: Abadi et al. 2016) + privacy ledger storing (query_id, ϵ _spent, composition).
 - **Example:** Suppose research query asks for aggregated counts. System computes noisy count: $\text{noisy} = \text{true_count} + N(0, \sigma^2)$, where σ computed from ϵ via Gaussian mechanism. The API returns $\text{noisy} + \text{epsilon_used}$.
-

6) Security: Key Flows & Escrow

- Keys per tenant + per purpose (ingest, pseudonymization, escrow).
 - HSM for root keys; KMS envelopes for per-tenant salts.
 - Decryption/de-pseudonymization flows must be signed and multi-party approved (e.g., 2-of-3 approvers) with time-boxed session tokens.
-

7) Deployment Patterns & Sizing (Kubernetes)

- **Namespace per tenant** for control-plane separation (not data). Use network policies to isolate.
 - **Pod types:** stateless workers (scale HPA), GPU pools (node pools) for OCR/ASR, stateful sets for Kafka & DBs (or managed services).
 - **Autoscaling triggers:** Kafka lag, CPU/GPU utilization, queue depth.
 - **Canary & Blue/Green** for policy bundles and models.
-

8) Observability: Concrete Metrics & Traces

- **Ingress:** requests_total, requests_latency_ms_bucket, auth_failures
 - **Processing:** messages_consumed, processing_latency_ms, detector_recall_estimate, false_negative_rate_est
 - **DP:** epsilon_spent_total, budgets_remaining
 - **Audit:** ledger_write_latency, ledger_anchor_latency
 - **Tracing:** request_id propagated (x-request-id) across services; spans for OCR, NER, policy_evaluate, redaction
-

9) Example: Near Miss Report (JSON)

```
{
  "document_id": "uuid",
  "detector_id": "ner.v2.1",
  "expected_entity": "SSN",
  "found": false,
```

```
"context": "Patient social security number appears in scanned image at 0x...",
"confidence_scores": [],
"ingest_ts": "2025-08-28T10:12:00Z",
"sample_for_labeling": "s3://bucket/sample/uuid.png"
}
```

10) Testing & Validation Plan

- Unit tests for parsers/policy bundles.
- Integration tests with synthetic PHI corpora (prescribed redaction expectations).
- Perf tests: k6 + custom Kafka load harness to validate 100k docs/s and p95 latency.
- Security: periodic pen-tests, automated SCA, CI gating for policy changes.

11) Implementation Roadmap (90 Days)

- **Week 0–2:** Bootstrapping infra: k8s, Kafka, API gateway, basic ingest API, schema registry.
- **Week 3–6:** Text/PDF path end-to-end: OCR, NER Tier-1, redaction, audit ledger v1, dashboards.
- **Week 7–9:** Audio path, streaming ASR & diarization, live redaction.
- **Week 10–12:** DICOM & structured data connectors (HL7/FHIR), entity graph & pseudonymization.
- **Weeks 13–16:** DP accountant, FL coordinator, adversarial robustness training, compliance validations.

12) Outstanding Decisions for Architect Sign-Off

1. On-prem vs SaaS defaults for top 10 customers.
2. Choice between Confluent vs Redpanda based on ops footprint and cost.
3. Level of commercial OCR/ASR to license vs build in-house.
4. Default ϵ/δ profiles to ship as safe presets.

13) Deliverables (next steps)

- Sequence diagrams in PlantUML for key flows.
- Terraform modules for infra bootstrap (k8s, kafka, object store).
- Prototype: a containerized pipeline that ingests PDF → OCR → Tier-1 NER → redacts → audit record.

End of Hour 3 deliverable.